

CCTEST

Testing and Repairing
Code Completion Systems



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE



Code Generation

- LLM-based code completion systems are commonly used as an advanced form of autocomplete.
 - A user begins typing code, and the LLM “infers” the rest of the code.
- These LLMs are trained on thousands of codebases
 - GitHub Copilot, for example, is trained from public repositories on GitHub.

```
3 class TrieNode:
4     def startsWith(self, prefix):
        if not prefix:
            return True
        if prefix[0] not in self.children:
            return False
        return self.children[prefix[0]].startsWith(prefix[1:])
```

Next (Alt+]) Previous (Alt+[) Accept (Tab) Open GitHub Copilot (Ctrl+Enter)

Problem

- Code that is functionally identical to a human programmer but with slight mutations to its content when used as input to an LLM can produce dramatically different results or even cause generation to fail.
 - As of the time of this paper, there is no automated way to test and improve these code completion systems.

```
def find_top_k(data_list, K):  
    length = len(data_list)  
    begin = 0  
    end = length - 1  
    index = divide(data_list, begin, end)
```

```
    if index == K:  
        return data_list[index]  
    elif index < K:  
        return find_top_k(data_list, K)  
    else:  
        return find_top_k(data_list, K)
```

```
def find_top_k(data_list, TOP_K):  
    length = len(data_list)  
    begin = 0  
    end = length - 1  
    index = divide(data_list, begin, end)
```

```
    if index == -1:  
        return -1  
    else:  
        return index
```


Solution: CCTest

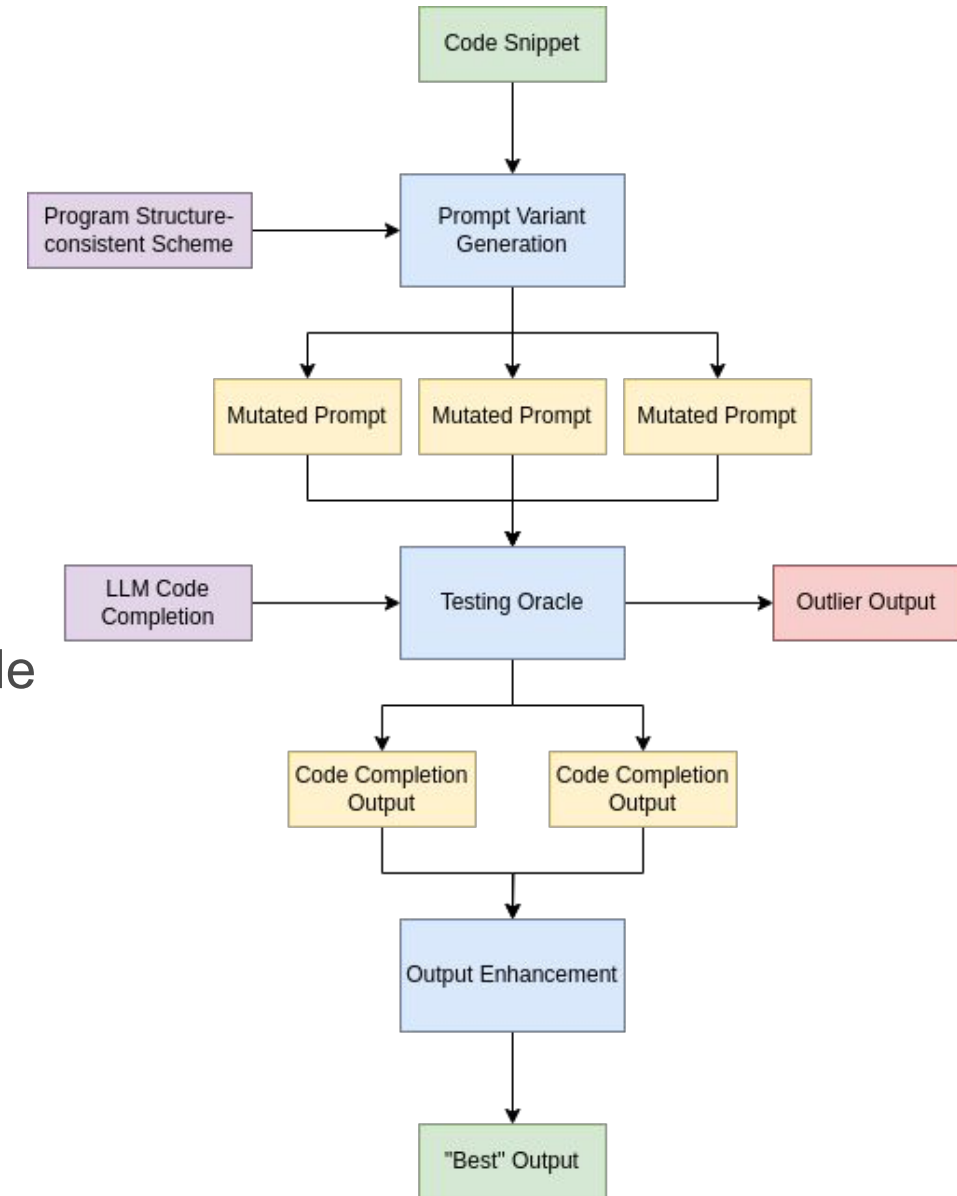
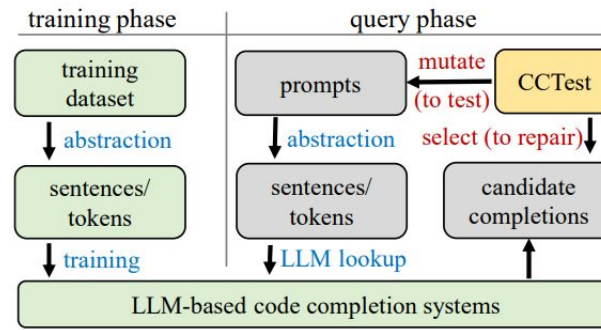
- Automated testing framework for LLM code completion systems.
- 3 main goals/steps:
 - Create mutated yet structurally similar input code prompts
 - Identify inconsistent outputs as potentially erroneous outliers
 - Enhance code generation outputs
- Tested on various code completion systems
 - GitHub Copilot, CodeParrot, GPT-Neo, GPT-J, and CodeGen
 - Completion system can be a “black box”
- Methods are applicable to any language
 - Only Python support currently



Approach/Methodology

Main Steps

- CCTest has 3 main steps
 - Generate mutants from given prompts (code snippets)
 - These variants should be identical functionally
 - These variants should be program structure-consistent
 - Identify outliers among the generated autocomplete code
 - Use metrics to find the average output appearance
 - Find outliers that deviate the furthest from that average
 - Label these as defective outputs / “bugs”
 - Exclude the outliers and recompute the average appearance output
 - Select the output closest to the average with outliers removed



Mutant Prompt Generation

- Prompt variant generation schemes:
 - REP_R, REP_C
 - Rename function parameters
 - REL_R, REL_C
 - Rename local variables
 - IRR
 - Replace arithmetic operators with semantically equivalent forms
 - RTF
 - Replace boolean expressions with semantically equivalent forms
 - GRA_R, GRA_C
 - Insert garbage code that does not alter program semantics
 - INI
 - Insert print statement into prompt

```
// seed prompt
def add(a,b):
    res = a + b
```

```
def add(a,Param1): // REP_R
    res = a + Param1
def add(a, Add_Param_b): // REP_C
    res = a + Add_Param_b
```

```
// seed prompt
def compare(a,b):
    res = a > b
```

```
def compare(a,b): // REL_R
    LocalVar1 = a > b
def compare(a,b): // REL_C
    compare_res = a > b
```

```
def addassign(a,b): // seed
    a += b
    return a
```

```
def addassign(a,b): // IRR
    a = a + b
    return a
```

```
def add (a,b,ignore): // seed
    if ignore:
        return a
```

```
def add (a,b,ignore): // RTF
    if ignore == (b==b):
        return a
```

```
// seed prompt
def add(a, b):
    res = a + b
```

```
def add(a, b): // GRA_R
    if (False): TempVar = a
    res = a+b
def add(a, b): // GRA_C
    if (b!=b): Add_TempVar = a
    res = a + b
```

```
// seed prompt
def add(a, b):
    res = a + b
```

```
def add(a, b): // INI
    print(b)
    res = a + b
```

Identifying Outlier Outputs

- For a given prompt and its mutant, iterate through every code completion output to find similarity between every other output
 - Similarity metric is based on Levenshtein edit distance
 - Obtains square matrix of similarity scores between every output
- If current output is less similar than another output's median similarity T times, then it is an outlier
 - T is a threshold hyperparameter


Algorithm 1 Outlier selection algorithm.

Input: \mathcal{O} : Code completion output set of size k
Input: T : threshold
Output: \mathcal{L} : Outliers

```
1:  $ScoreMatrix = []$ 
2: for  $i$  in 1 to  $k$  do
3:   for  $j$  in  $i$  to  $k$  do  $ScoreMatrix[i][j] = Sim(o_i, o_j)$ 
4:  $Normalize(ScoreMatrix)$ 
5: for  $i$  in 1 to  $k$  do
6:    $count = 0$ 
7:   for  $j$  in 1 to  $k$  do
8:     if  $ScoreMatrix[i][j] < Median(ScoreMatrix)$  then
9:        $count = count + 1$ 
10:    if  $count \geq T$  then
11:       $\mathcal{L}.append(o_i)$ 
12:    break
13: return  $\mathcal{L}$ 
```


Code Completion System Enhancement

- Find average similarity score of all code completion outputs
 - Excludes previously found outliers
 - Similarity metric based on Levenshtein edit distance
 - Essentially average value in previously found similarity square matrix (excluding outliers)
- Output with similarity score closest to this average is chosen as the “enhanced” code completion output



Other Experimental Setup Info

- Dataset used for prompts are from LeetCode and CodesearchNet
 - Total of 2,910 programs + 19,898 mutant prompts = 22,808 total prompts
- Parse Python code into concrete syntax tree and determine which mutations can be performed
- Evaluated code completion systems:
 - GitHub Copilot
 - CodeParrot
 - GPT-Neo
 - GPT-J
 - CodeGen



LeetCode

Findings/Results

Mutant Prompt Generation

- pycode-similar used to determine distance between prompts and mutants and, therefore, effectiveness of CCTEST prompt mutation
 - Different distance than Levenshtein edit distance used for outputs
 - pycode-similar based on AST structures
- 98.46% of mutated prompts had distance < 0.1
 - Authors conclude that CCTEST successfully generates structurally-consistent prompt mutants

TABLE IV
DISTRIBUTION OF STRUCTURAL CONSISTENCY SCORES.

Distance	[0, 0.05]	[0.05, 0.1]	[0.1, 0.15]	[0.15, 0.2]	[0.2, 0.9]	[0.9, 1.0]
Freq. (%)	90.16	8.30	0.92	0.30	0.32	0
Cumulative Freq. (%)	90.16	98.46	99.38	99.68	100.0	100.0

Identifying Outlier Outputs

- Significant number of outliers=defects found for every value of T
 - Copilot tended to produce the fewest defects
- Increasing threshold T decreases number of detected outliers
- Authors manually validated whether or not detected outliers were correctly classified as “defects”
 - Obtained TPs, FNs, precision, recall, F1 score
 - Determined T=9 appears to be best threshold hyperparameter, F1=0.865
- All mutations appear to contribute towards defects pretty equally

TABLE VI
ASSESSING CCTEST’S FINDINGS WITH MANUAL INVESTIGATION.

	T=1	T=3	T=5	T=7	T=9
TP	216	342	429	537	689
FN	26	38	70	90	104
Precision	0.270	0.427	0.536	0.671	0.861
Recall	0.892	0.900	0.859	0.856	0.868
F1 score	0.414	0.579	0.660	0.752	0.865

TABLE V
OVERVIEW OF OUTLIER DETECTION RESULTS. FOR EACH SYSTEM, WE USE 4909 LEETCODE PROMPTS AND 17899 CODESEARCHNET PROMPTS TO TEST.

System	#No Results	#Outliers				
		T= 1	T= 3	T= 5	T= 7	T= 9
Copilot	4+37	3003 + 12101	1347 + 7928	803 + 5570	559 + 3899	293 + 2184
CodeParrot	1+0	4798 + 17605	4379 + 16118	3631 + 13368	2359 + 9023	904 + 3778
CodeParrot-small	2+0	4812 + 17611	4469 + 16069	3776 + 13454	2470 + 9344	1033 + 4009
GPT-J	0+0	4606 + 17280	3776 + 15073	2832 + 12005	1786 + 7875	586 + 3174
GPT-NEO-13B	1+0	4729 + 17427	4088 + 15495	3311 + 12536	2120 + 8462	794 + 3463
GPT-NEO-125M	0+2	4734 + 17509	4281 + 15556	3654 + 12907	2523 + 8966	1079 + 3700
Codegen-2B-mono	2+9	4661 + 17221	3794 + 15112	2731 + 12241	1638 + 8460	639 + 3761
Codegen-6B-mono	0+0	4578 + 17016	3575 + 14595	2493 + 11546	1452 + 7866	584 + 3559
Total	10+48	35921 + 133770	29709 + 115946	23231 + 93627	14907 + 63895	5912 + 27628

Code Completion System Enhancement

- Found BLEU score and Levenshtein edit similarity between code completion outputs and ground truth from datasets
 - Calculated enhancement ratio as $(r = (s' - s) / s)$, where s' is score after enhancement and s is score before enhancement (no use of mutant prompts)
 - Found average enhancement to be 40% by BLEU score and 62% to 73% by Levenshtein edit distance, depending on the evaluated dataset
- Found each mutation contributed pretty equally to enhancement
- Performed human study in which experts were tasked to score random code completions from 1-5
 - Some were enhanced, some unenhanced
 - Average unenhanced score was 2.32 while enhanced was 3.16
 - 97.8% of respondents believed that enhanced outputs looked better than or equivalent to unenhanced outputs

Conclusion

- CCTEST is successful at:
 - Creating program structure-consistent code prompt mutations
 - Finding defects in LLM code completion systems
 - Enhancing LLM code completion outputs
- CCTEST would likely be cost-effective to deploy, especially on the server side (e.g. by GitHub Copilot)
- Authors hope CCTEST sets the tone for future work in testing and improvement of code completion systems

Paper Reference

```
@inproceedings{li2023cctest,  
  title={Cctest: Testing and repairing code completion systems},  
  author={Li, Zongjie and Wang, Chaozheng and Liu, Zhibo and Wang, Haoxuan and Chen, Dong and Wang, Shuai and Gao, Cuiyun},  
  booktitle={2023 IEEE/ACM 45th International Conference on Software Engineering (ICSE)},  
  pages={1238--1250},  
  year={2023},  
  organization={IEEE}  
}
```